

Verification on CAE


By Abhinav Nandwani

Use this guide to compile a SystemVerilog test, inspect its source and waveforms, generate coverage, and prove that the test catches a known bug. The terminal and GUI sections use the same small register example. The final section connects this workflow to the team's test plan and verification scaffolding.

Sign in to Guacamole

Guacamole displays a CAE Linux desktop in your browser. The tools run on the CAE machine. Use your own UW NetID and confirm that your account can reach CAE Linux services before the lab.

1. On your laptop, open <https://guacamole.cae.wisc.edu> in a browser.
2. Complete the UW web sign-in with your NetID and password, then complete MFA if requested.
3. At the Linux login screen, enter your UW NetID and password again. Use your NetID credentials at this second prompt too. The old separate CAE username and password are not the instructions for this service.
4. Wait for the Linux desktop. Open Apps in the upper-left corner and choose Terminal under Favorites. Maximize the terminal by double-clicking its title bar.
5. Keep the Guacamole tab open. A terminal inside this desktop is already on CAE; do not SSH back into CAE from it.



Open Terminal from the CAE Apps menu.

If you already have an active session, the browser may reconnect without showing every login screen. If sign-in fails, record which prompt failed: UW web sign-in, MFA, or the Linux desktop login. These are different stages.

Official CAE login instructions and screenshots: <https://kb.wisc.edu/cae/163323>

Enter the Synopsys environment

Run these commands in the CAE terminal, one line at a time. Do not type a prompt such as \$ or synopsys> before a command.

```
module load synopsys/suite
synopsys-run
command -v vcs dc_shell verdi icc2_shell pt_shell
```

The prompt becomes synopsys>. The module selects the site setup; the container provides the operating environment expected by the tools. Load the module before entering the container and repeat this setup for every new terminal used for Synopsys tools.

```
synopsys> command -v vcs dc_shell verdi icc2_shell pt_shell
/srv/auto/apps/synopsys-2026/vcs/Y-2026.03/bin/vcs
/srv/auto/apps/synopsys-2026/syn/Y-2026.03/bin/dc_shell
/srv/auto/apps/synopsys-2026/verdi/Y-2026.03/bin/verdi
/srv/auto/apps/synopsys-2026/icc2/X-2025.06-SP3/bin/icc2_shell
/srv/auto/apps/synopsys-2026/prime/Y-2026.03/bin/pt_shell
synopsys>
```

Tool paths inside the working CAE Synopsys container.

Companion code

RTL: https://github.com/abhinavnandwani/cae-synopsys-guides/blob/main/lab/rtl/sb_flop.v

Testbench: <https://github.com/abhinavnandwani/cae-synopsys-guides/blob/main/lab/tb/tb.sv>

Simulation runner: https://github.com/abhinavnandwani/cae-synopsys-guides/blob/main/lab/run_lab.py

Commands below use ~/cae-synopsys-guides/lab as the exercise directory. The recorded screenshots show the original dated demonstration folder.

Work confidently in the terminal

The current directory affects relative paths. `pwd` prints it, `ls` lists files, `cd ..` moves to the parent, and `cd ~/cae-synopsys-guides/lab` returns to the exercise. A leading `/` means an absolute path; `~` means your home directory. Quote paths containing spaces.

Prompt or location	Commands that belong there
Laptop Terminal or PowerShell	<code>ssh</code> to reach the CAE host shell.
CAE host shell	<code>module load synopsys/suite</code> , then <code>synopsys-run</code> .
Container <code>synopsys></code>	Linux commands, <code>python3 run_lab.py</code> , <code>vcs</code> , <code>verdi</code> , <code>dc_shell</code> , and <code>icc2_shell</code> .
Tool prompt such as <code>dc_shell></code>	Tool Tcl commands such as <code>read_ddc</code> , <code>report_timing</code> , and <code>help</code> .

Use these from the exercise folder to inspect source and logs without changing them:

```
ls -lh
cat rtl/sb_flop.v
sed -n '1,80p' tb/tb.sv
less runs/YOUR_RUN/compile.log
grep -nE 'Error|Fatal|Warning' runs/YOUR_RUN/compile.log
tail -n 30 runs/YOUR_RUN/simulate.log
```

Replace `YOUR_RUN` with an actual directory printed by the runner. In `less`, use Space to advance, `/Error` then Enter to search, `n` for the next match, and `q` to return to the shell. An Up-arrow recalls a command; Tab completes a path. Read a command before rerunning it.

For an interactive foreground command that is stuck, `Ctrl+C` requests interruption. For a GUI launched with `&`, close the application through its File menu. `jobs` lists jobs started by that shell. Keep logs: `command > run.log 2>&1` sends both standard output and errors to a file. Immediately after a command, `echo $?` reports its exit status, but a zero status alone does not prove the design passed.

Connect without the browser when useful

For terminal-only work, open a terminal on your laptop and run `ssh YOUR_NETID@best-tux.cae.wisc.edu`, then load the module and enter the container as above. Use Guacamole for the GUI steps in this handout. SSH and Guacamole may reach different hosts, but your CAE home directory is shared.

When finished, save work, close the EDA applications, and run `exit` to leave the container. Log out of the Linux desktop from the upper-right system menu. Closing the browser alone can leave the session running during CAE's two-hour reconnection window.

Run the verification exercise from the terminal

Inside the Synopsys container, start from the companion folder:

```
cd ~/cae-synopsys-guides/lab
python3 run_lab.py simulation
```

The runner creates a new directory under `runs/` for every invocation. It compiles the design, executes a passing simulation, builds a coverage report, and executes a separate intentionally failing simulation. The final `LAB_RESULT=PASS` means these checks behaved as expected, including detection of the deliberate failure.

```
synopsys> grep -E 'UVM_INFO|SB_SNPS_SIM_PASS|$finish at' simulate.log
UVM_INFO /srv/auto/apps/synopsys-2026/vcs/Y-2026.03/etc/uvm-1.2/base/uvm_root.svh(407
) @ 0: reporter [UVM/RELNOTES]
UVM_INFO /userspace/nandwani2/siliconbadgers-cae-lab-20260923/tb/tb.sv(11) @ 0: repor
ter [ACCESS] UVM package and macros are available
SB_SNPS_SIM_PASS
synopsys>
```

The saved passing log contains the UVM access message and `SB_SNPS_SIM_PASS`.

File in the printed run directory	How to use it
<code>compile.log</code>	Start here for syntax, missing include, package, and elaboration errors.
<code>simulate.log</code>	Inspect test diagnostics, simulation time, and the pass marker.
<code>simv</code>	The compiled simulation executable for this build.
<code>smoke.fsdb</code>	Waveform data for the passing run.
<code>simv.vdb</code>	Collected coverage data.
<code>coverage_report/</code>	URG's HTML coverage report.
<code>negative/expected_failure.log</code>	Evidence that the injected fault was detected.
<code>result.json</code>	Exact commands, working directories, elapsed time, raw exit codes, and verdict.

Change into the exact run directory the runner printed. Do not guess that an older directory is the latest run. Inspect `result.json` and the logs even when the final result says `PASS`. A working test harness and a correct design are separate questions.

Understand and reproduce the terminal commands

The example has `rtl/sb_flop.v` and `tb/tb.sv`. The testbench instantiates the register, drives a 10 ns clock, checks reset, checks captured data, and writes an FSDB. UVM 1.2 is imported to check that the installed integration works; this example is a procedural testbench, not a complete class-based UVM environment.

To run the commands directly, create a separate manual run directory first:

```
SB_LAB="$HOME/cae-synopsys-guides/lab"
SB_RUN=$(mktemp -d "$SB_LAB/runs/manual-sim-XXXXXX")
cd "$SB_RUN"
export VERDI_HOME="$(dirname "$(dirname "$(command -v verdi)"))"
vcs -full64 -sverilog -timescale=1ns/1ps \
  -ntb_opts uvm-1.2 -debug_access+all -kdb \
  -cm line+cond+tgl+branch+assert \
  "$SB_LAB/rtl/sb_flop.v" "$SB_LAB/tb/tb.sv" \
  -top tb -o simv > compile.log 2>&1
echo $?
./simv -cm line+cond+tgl+branch+assert > simulate.log 2>&1
grep -nE 'Error|Fatal|SB_SNPS_SIM_PASS' simulate.log
urg -full64 -dir simv.vdb -report coverage_report \
  > coverage.log 2>&1
```

Stop after compilation if it fails. The shell does not automatically stop this sequence for you. The supplied Python runner is preferable for repeatable runs because it checks diagnostics and required outputs and gives each command a 180-second limit.

`-full64` selects the working 64-bit VCS mode on this installation. `-sverilog` enables SystemVerilog parsing. `-top tb` selects the testbench top. `-debug_access+all` `-kdb` retains information needed for source-level debugging. The `-cm` options enable the listed coverage types; they do not create missing assertions or functional coverage goals.

Inspect the test before trusting it


Read the checks in `tb/tb.sv`. At 11 ns the test checks the reset value and drives one onto `d`. It checks `q` at 21 ns, then drives zero. It checks again at 31 ns. These checks occur after the relevant rising edge so the nonblocking assignment has completed. A real interface test should define its sampling convention just as explicitly.

Open Verdi and navigate the design

From the passing simulation directory, launch the GUI inside Guacamole:

```
verdi -ssf smoke.fsd
```

Keep that terminal available for launch diagnostics. Maximize Verdi by double-clicking its title bar. The hierarchy pane identifies instances; the source pane shows their code; the lower nWave pane displays recorded signal values.




The passing test open in Verdi with its source and waveform panes.

1. In the hierarchy pane, select `tb`. Expand it and select `dut` to inspect the instantiated `sb_flop` module. Confirm that you opened the intended source and top.
2. Return to `tb` to read the stimulus and checks alongside the RTL. Locate the clock generator, reset release, assignments to `d`, and `$fatal` checks.
3. In the lower nWave pane, choose Signal, then Get Signals. Select the `/tb` scope in the dialog.
4. Select `clk`, `rst`, `d`, and `q`. In the tested interface, selecting a signal and clicking Apply adds it to the waveform list. Repeat for the four signals, then click OK.
5. In nWave, choose View, then Zoom, then Zoom All. This displays the entire saved simulation rather than just its initial few nanoseconds.

If the source view is empty, first confirm that you launched from the run directory that contains the matching build data. Do not mix one build's debug database with another build's waveform. If the waveform list is empty, adding signals is still required even though the FSDB loaded successfully.

Debug the waveform rather than just displaying it

Start with all four signals visible and the full 0 to 31 ns interval. The recorded view uses picoseconds on the ruler: 10000 ps is 10 ns. Signal values in the value column correspond to the current cursor position, not necessarily the end of the run.



Clock, reset, data, and output shown over the complete saved test.

1. Click in the waveform near the first rising edge of `clk`. Follow the same vertical time position through `rst`, `d`, and `q`.
2. Move the time cursor before and after the edge. Before the first rising edge, `q` is unknown. At 5 ns, the asserted synchronous reset sets it to zero.
3. Inspect the reset release and data change at 11 ns. The output should not update immediately because the register samples on the next rising edge.
4. Inspect 15 ns. `q` captures one. Inspect 21 ns when `d` returns to zero, then 25 ns when `q` captures zero.
5. Use the nWave View, Zoom menu to inspect a smaller interval, then Zoom All to recover the overview. Keep the clock visible when investigating a data transition.

Time	Expected observation
Before 5 ns	<code>q</code> is unknown; no active clock edge has applied reset yet.
5 ns	<code>q</code> becomes zero while <code>rst</code> is high.
11 ns	<code>rst</code> becomes zero and <code>d</code> becomes one.
15 ns	<code>q</code> becomes one.
21 ns	<code>d</code> becomes zero.
25 ns	<code>q</code> becomes zero.
31 ns	Final check completes and simulation ends.

Write down the first time observed behavior differs from expected behavior. Then return to the source driving or sampling that signal. A waveform file existing proves only that recording worked; the test's checks and these transitions establish the example's behavior.

Save the signal arrangement with nWave File, Save Signal. Give the file a descriptive `.rc` name in the run directory. Use File, Restore Signal to load that arrangement in a later session with the matching FSDB. This saves a view configuration, not new simulation results. For readable screenshots, enlarge Monospaced Font under Tools, Preferences, General, Appearance.

Investigate the deliberately failing run

The runner invokes the simulation again with `+INJECT_FAILURE` in a separate `negative/` directory. At 11 ns, the test forces the register output to zero. The expected one is therefore missing at the 21 ns check.

```
cat negative/expected_failure.log
verdi -ssf negative/smoke.fsdb
```

```
synopsys> tail -n 10 negative/expected_failure.log
*Verdi* : Create FSDB file 'smoke.fsdb'
*Verdi* : Begin traversing the scope (tb), layer (0).
*Verdi* : End of traversing.
UVM_INFO /userspace/nandwani2/siliconbadgers-cae-lab-20260923/tb/tb.sv(11) @ 0: reporter [ACCESS] UVM package and macros are available
Fatal: "/userspace/nandwani2/siliconbadgers-cae-lab-20260923/tb/tb.sv", 17: tb: at time 21000 ps
data failed
$finish called from file "/userspace/nandwani2/siliconbadgers-cae-lab-20260923/tb/tb.sv", line 17.
$finish at simulation time                21000
      V C S   S i m u l a t i o n   R e p o r t
Time: 21000 ps
synopsys> █
```

The captured failing run stops at 21000 ps with the data failed diagnostic.

In the failing waveform, add the same four signals. Confirm that `d` becomes one but `q` stays zero, then locate the failed check at 21 ns in the source. Compare against the passing run at 15 ns and 21 ns. A separate viewer is useful for comparing the two results; close each viewer when finished.

On the tested installation, VCS returned raw exit code zero even for this fatal diagnostic. The runner therefore requires the positive pass marker, rejects unexpected error or fatal diagnostics, and separately verifies that the deliberate failure reports `data failed` without the pass marker. Do not use exit zero as the only test verdict.

Preserve both sides of the comparison

Keep the passing and failing logs and waveforms in their separate directories. Record the expected failure text and time. When adapting this pattern, inject a fault that violates a specific requirement and confirm that the intended checker catches it. A crash during compilation is not evidence that a behavioral checker works.

Inspect coverage in the GUI

In the CAE desktop file manager, open the passing run's coverage_report folder and open dashboard.html with Firefox. Alternatively, in a CAE host terminal, run `firefox /absolute/path/to/coverage_report/dashboard.html`. This browser runs on CAE; the file is not on your laptop.

Module Definition
dashboard | hierarchy | modlist | groups | tests | asserts

Line Toggle Branch

Module : **sb_flop**

SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT
95.83	100.00		87.50	100.00	

Source File(s) :
/userspace/nandwani2/siliconbadgers-cae-lab-20260923/rtl/sb_flop.v

Module self-instances :

NAME	SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT
tb_dut	95.83	100.00		87.50	100.00	

Module Instance : **tb_dut**

Instance :

SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT
95.83	100.00		87.50	100.00	

Instance's subtree :

SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT
95.83	100.00		87.50	100.00	

Parent :

SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT	NAME
73.83	84.00		87.50	50.00		tb

Subtrees :

NAME	SCORE	LINE	COND	TOGGLE	BRANCH	ASSERT
no children						

Since this is the module's only instance, the coverage report is the same as for the module.

Line Coverage for Module : sb_flop

Line No.	Total	Covered	Percent
TOTAL	3	3	100.00
ALWAYS	3	3	100.00

```

2          always @(posedge clk)
3              if (rst) q <= 1'b0;
4              else   q <= d;
  
```

Toggle Coverage for Module : sb_flop

Total	Covered	Percent	
Totals	4	3	75.00
Total Bits	8	7	87.50
Total Bits 0->1	4	3	75.00
Total Bits 1->0	4	4	100.00
Ports	4	3	75.00
Port Bits	8	7	87.50
Port Bits 0->1	4	3	75.00
Port Bits 1->0	4	4	100.00

Port Details

Name	Toggle	Toggle 1->0	Toggle 0->1	Direction
clk	Yes	Yes	Yes	INPUT
rst	No	Yes	No	INPUT
d	Yes	Yes	Yes	INPUT
q	Yes	Yes	Yes	OUTPUT

Branch Coverage for Module : sb_flop

Line No.	Total	Covered	Percent
Branches	2	2	100.00
IF	3	2	100.00

URG's DUT page shows source coverage and the missing reset toggle.

1. On the dashboard, check the run date, tool version, command line, and test count. This example report contains one passing test; the deliberately failing run is separate.
2. Click hierarchy at the top. Expand tb if necessary, then click its dut child. The module page should identify sb_flop and the instance tb_dut.
3. Inspect Line and Branch. The tested DUT shows 100% for both. Match the covered source statements to the reset and data paths in the RTL.
4. Inspect Toggle and its Port Details table. The DUT reports 87.5% toggle coverage. rst has a falling transition but no rising transition because the test starts in reset and never reasserts it.
5. Compare the DUT with the dashboard totals. The totals also include testbench and imported UVM code. Always record the measured scope and metric with a percentage.

As a follow-on exercise, reassert reset after capturing data, check the resulting output at the appropriate clock edge, and regenerate the report in a fresh run. Confirm both the new behavior and the previously missing transition. Increasing coverage without a meaningful check is not sufficient.

Condition and assertion coverage do not have meaningful DUT targets in this tiny example. High code coverage does not establish arithmetic correctness, protocol behavior, corner cases, or interactions between units. The project test plan must connect requirements to stimulus, checks, and suitable coverage goals.

Turn the exercise into verification infrastructure

For each project layer, specify the behavior to verify, reference model or expected result, stimulus, checks, coverage, and reproducible command. Decide where C++, SystemVerilog, UVM, or another method fits the layer. The availability of a UVM package does not by itself choose the methodology.

What a runnable test should retain	Why it is needed
Source revision and source manifest	Identifies the exact implementation and test compiled.
Tool version and full commands	Makes setup and build differences diagnosable.
Seed and configuration when applicable	Reproduces randomized or configurable behavior.
Timeout and explicit verdict	Distinguishes a hang, a crash, and a failed check.
Logs, waveform, and coverage	Preserves evidence for debugging and review.
Known failing case	Shows that the checker rejects incorrect behavior.

Troubleshoot in a consistent order

If compilation fails, read the first relevant error before investigating later errors that may be consequences. Check the top module, source list, package and include order, and SystemVerilog flags. If simulation starts but no pass marker appears, inspect where it stops and whether the expected checker ran. If the GUI is blank, check the selected FSDB, scope, signal list, and matching build directory.

If a tool prints only a CAE launcher warning, return to the host shell and load the module before entering the container. For a license failure, retain the exact diagnostic and tool version. For an unresponsive GUI, first allow startup or loading to complete; do not repeatedly start duplicate copies that each consume resources.

Completion check

Complete the passing and deliberately failing runs, explain the waveform at each expected edge, identify at least one coverage limitation, and show another teammate how to reproduce the run from its command and source revision. The team's project deliverables remain the architecture-based test plan, per-layer methodology proposal, and hardened runnable verification scaffolding.

Tested on CAE with VCS Y-2026.03_Full64, Verdi Y-2026.03, and URG Y-2026.03. Setup reference: <https://kb.wisc.edu/cae-software-guide>

Verdi overview: <https://www.synopsys.com/verification/debug/verdi.html>